
Accelerating CNN Inference Using SIMD Techniques

Group 40

312553049 單綿恆 312581031 張清華 31351007 陳晶

Outline

1. Abstract

2. Method

3. Experiments

Outline

1. Abstract

- Motivation
- AlexNet

2. Method

3. Experiments

Outline

1. Abstract

- Motivation
- AlexNet

2. Method

3. Experiments

Abstract – Motivation

1. CNNs are **widely used** in various scenarios, with AlexNet being one of the most representative examples.
2. Small **IoT devices lack GPU** capabilities.
3. Dive into SIMD instructions for low-level optimization.
4. Apply SIMD for AI acceleration and inference optimization.

Outline

1. Abstract

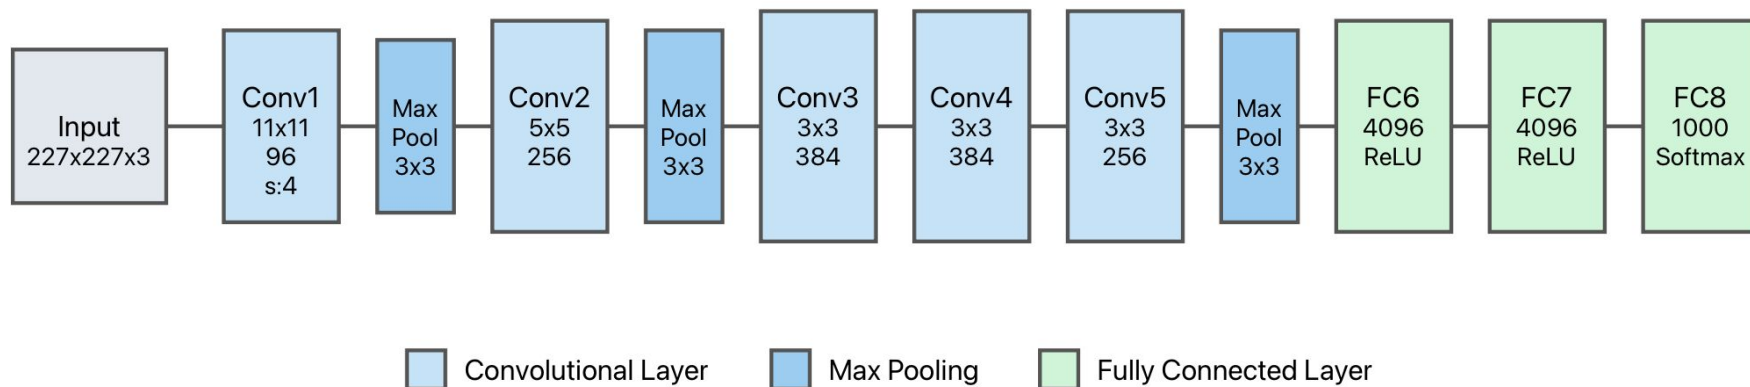
- Motivation
- AlexNet

2. Method

3. Experiments


Abstract – AlexNet

Architecture



Abstract – AlexNet

Time Consumption Analysis



Outline

1. Abstract

2. Method

- Hardware
- Software

3. Experiments

Outline

1. Abstract

2. Method

- Hardware
- Software

3. Experiments

Method – Hardware

Adding SIMD ISA Extension

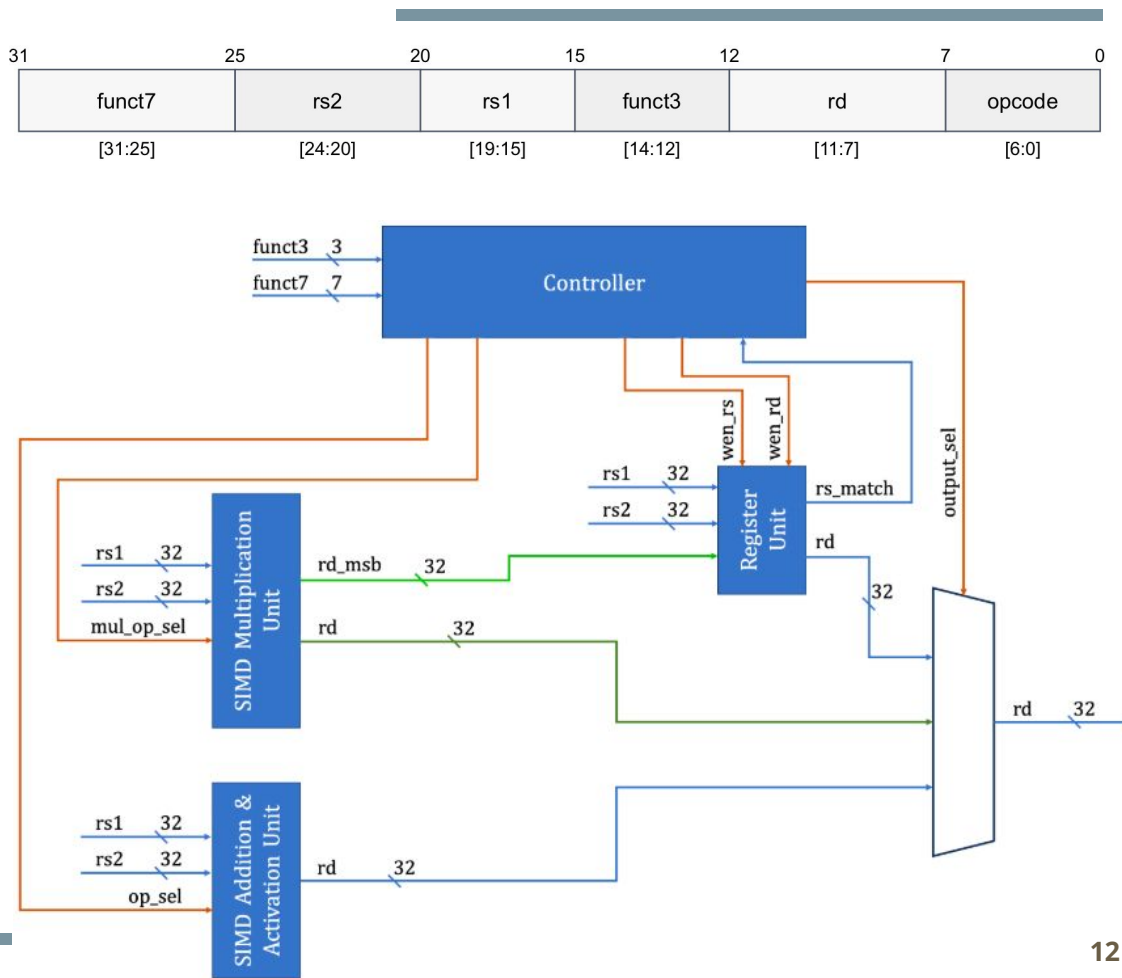
Four instruction type:

- Add
- Subtract
- Multiply
- Quantize

Type	Vector-Vector	Vector-Scalar
ADD	sADDI8I8S.vv	sADDI8I8S.vx
	sADDI16I16S.vv	sADDI16I16S.vx
SUB	sSUBI8I8S.vv	sSUBI8I8S.vx
	sSUBI16I16S.vv	sSUBI16I16S.vx
PMUL	sPMULI8I16S.vv.L	sPMULI8I16S.vx.L
	sPMULI8I16S.vv.H	sPMULI8I16S.vx.H
AMUL	sAMULI8I8S.vv.NQ	sAMULI8I8S.vx.NQ
	sAMULI8I8S.vv.AQ	sAMULI8I8S.vx.AQ
QNT	sQNTI16I8S.vv.NQ	
	sQNTI16I8S.vv.AQ	

Method – Hardware

Custom Function Unit — SIMD Execution Engine




Method – Hardware



	31	25	20	15	12	7	0					
	funct7		rs2		rs1		funct3		rd		opcode	
	[31:25]		[24:20]		[19:15]		[14:12]		[11:7]		[6:0]	
s ADD i8i8S.vx	7b1000000						3b000				SIMD	
s ADD i16i16S.vx	7b1000000						3b001				SIMD	



+



||




+


+

+

+



||



Method – Hardware

ALU Implementation

```
// 8-bit wire assignment
for (i <- 0 until 4) {
  rs1ByteArray(i) := io.rs1(8 * i + 7, 8 * i)
  rs2ByteArray(i) := io.rs2(8 * i + 7, 8 * i)

  rdByteArray(i) := MuxLookup(
    io.opSel.asUInt,
    DontCare,
    Seq(
      AddSubActivationOp.ADDI8I8S_VV.asUInt -> (rs1ByteArray(i).asSInt + rs2ByteArray(i).asSInt).asUInt,
      AddSubActivationOp.SUBI8I8S_VV.asUInt -> (rs1ByteArray(i).asSInt - rs2ByteArray(i).asSInt).asUInt,
      AddSubActivationOp.ADDI8I8S_VX.asUInt -> (rs1ByteArray(i).asSInt + rs2ByteArray(0).asSInt).asUInt,
      AddSubActivationOp.SUBI8I8S_VX.asUInt -> (rs1ByteArray(i).asSInt - rs2ByteArray(0).asSInt).asUInt,
      AddSubActivationOp.SCMPL8.asUInt -> Mux(
        rs1ByteArray(i).asSInt <= rs2ByteArray(i).asSInt,
        rs2ByteArray(i),
        rs1ByteArray(i)
      )
    )
  )
}
```

Outline

1. Abstract

2. Method

- Hardware
- **Software**

3. Experiments

Method – Software

Convolution Implementation

```
void Conv::execPerLayerNaiveQuant() {
    int16_t tempINT16_Buffer[100000] = {0};
    int16_t temp_C[4] = {0}, temp_D[4] = {0};
    int8_t temp_A[4] = {0}, temp_B[4] = {0}, temp_E[4] = {0};

    for (int n = 0; n < info->kernel.N; n++) {
        for (int oh = 0; oh < output->H; oh++) {
            for (int ow = 0; ow < output->W; ow++) {
                output->data[n * output->H * output->W + oh * output->W + ow] = info->bias.data[n];
                for (int c = 0; c < info->kernel.C; c++) {
                    for (int kh = 0; kh < info->kernel.H; kh++) {
                        for (int kw = 0; kw < info->kernel.W; kw += 4) {
                            // ... (rest of the code is cut off)
                        }
                    }
                }
            }
        }
    }
}
```


Method – Software

Convolution Implementation

temp_A

A11	A12	A13	A14
-----	-----	-----	-----

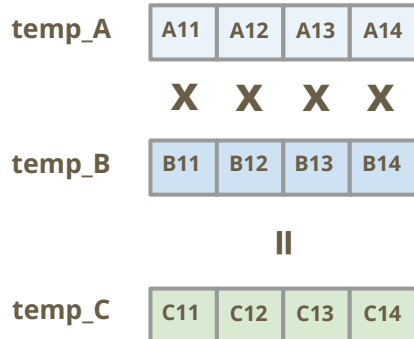
temp_B

B11	B12	B13	B14
-----	-----	-----	-----

```
for (int kw = 0; kw < info->kernel.W; kw += 4) {
    for (int i = 0; i < 4; i++) {
        if (kw + i >= info->kernel.W) {
            temp_A[i] = 0;
            temp_B[i] = 0;
        } else {
            temp_A[i] = input->data[c * input->H * input->W
                                + (oh + kh) * input->W
                                + (ow + kw + i)];
            temp_B[i] = info->kernel.data[(n * info->kernel.C + c) * info->kernel.H + kh] * info->kernel.W
                                + kw
                                + i];
        }
    }
    sPMULI8I16S_vv_L(temp_C, temp_A, temp_B);
    sPMULI8I16S_vv_H(temp_C + 2, temp_A, temp_B);
    for (int i = 0; i < 4; i++)
        tempINT16_Buffer[n * output->H * output->W + oh * output->W + ow] += temp_C[i];
}
```

Method – Software


Convolution Implementation



```
for (int kw = 0; kw < info->kernel.W; kw+=4) {
    for (int i = 0; i < 4; i++) {
        if (kw + i >= info->kernel.W) {
            temp_A[i] = 0;
            temp_B[i] = 0;
        }else{
            temp_A[i] = input->data[c * input->H * input->W
                                + (oh + kh) * input->W
                                + (ow + kw + i)];
            temp_B[i] = info->kernel.data[((n * info->kernel.C + c) * info->kernel.H + kh) * info->kernel.W
                                + kw
                                + i];
        }
    }
    SPMULI8I16S_vv_L(temp_C, temp_A, temp_B);
    SPMULI8I16S_vv_H(temp_C+2, temp_A, temp_B);
    for (int i = 0; i < 4; i++)
        tempINT16_Buffer[n * output->H * output->W + oh * output->W + ow] += temp_C[i];
}
```

Method – Software

GEMM Implementation

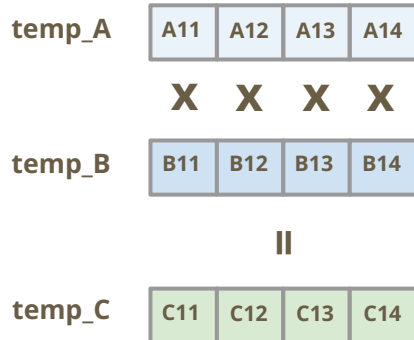


```
void Gemm::execPerLayerAdvanceQuant() {
    int16_t tempINT16_Buffer[100000] = {0};
    int index_A, index_B, index_C;
    int8_t temp_A[4] = {0}, temp_B[4] = {0};
    int16_t temp_C[4] = {0}, temp_D[4] = {0};

    for (int m = 0; m < input->H; m++) {
        index_A = m * input->W; // M * K
        index_C = m * output->W; // M * N
        for (int k = 0; k < input->W; k++) {
            index_B = k * info->weight.W; // K * N
            for (int n = 0; n < info->weight.W; n += 4) {
                for (int i = 0; i < 4; ++i) {
                    if (n + i >= info->weight.W)
                        temp_A[i] = 0;
                    else
                        temp_A[i] = input->data[index_A + k];
                }
                *(int32_t*)temp_B = *(int32_t*)&(info->weight.data)[index_B + n];
                int output_index = index_C + n;
                sPMULI8I16S_vv_L(temp_C, temp_A, temp_B);
                sPMULI8I16S_vv_H(temp_C + 2, temp_A, temp_B);
                sADDI16I16S_vv(tempINT16_Buffer + output_index,
                               tempINT16_Buffer + output_index,
                               temp_C);
                sADDI16I16S_vv(tempINT16_Buffer + output_index + 2,
                               tempINT16_Buffer + output_index + 2,
                               temp_C + 2);
            }
        }
        for (int n = 0; n < info->weight.W; n++)
            tempINT16_Buffer[index_C + n] += info->bias.data[index_C + n];
    }
}
```

Method – Software

GEMM Implementation



```
void Gemm::execPerLayerAdvanceQuant() {
    int16_t tempINT16_Buffer[100000] = {0};
    int index_A, index_B, index_C;
    int8_t temp_A[4] = {0}, temp_B[4] = {0};
    int16_t temp_C[4] = {0}, temp_D[4] = {0};

    for (int m = 0; m < input->H; m++) {
        index_A = m * input->W; // M * K
        index_C = m * output->W; // M * N
        for (int k = 0; k < input->W; k++) {
            index_B = k * info->weight.W; // K * N
            for (int n = 0; n < info->weight.W; n += 4) {
                for (int i = 0; i < 4; ++i) {
                    if (n + i >= info->weight.W)
                        temp_A[i] = 0;
                    else
                        temp_A[i] = input->data[index_A + k];
                }
                *(int32_t*)temp_B = *(int32_t*)&(info->weight.data)[index_B + n];
                int output_index = index_C + n;
                SPMULI8I16S_vv_L(temp_C, temp_A, temp_B);
                SPMULI8I16S_vv_H(temp_C + 2, temp_A, temp_B);
                sADDI16I16S_vv(tempINT16_Buffer + output_index,
                               tempINT16_Buffer + output_index,
                               temp_C);
                sADDI16I16S_vv(tempINT16_Buffer + output_index + 2,
                               tempINT16_Buffer + output_index + 2,
                               temp_C + 2);
            }
        }
    }
    for (int n = 0; n < info->weight.W; n++)
        tempINT16_Buffer[index_C + n] += info->bias.data[index_C + n];
}
```

Outline

1. Abstract

2. Method

3. Experiments

- Environment
- Serial vs SIMD
- Summary

Outline

1. Abstract

2. Method

3. Experiments

- Environment
- Serial vs SIMD
- Summary

Experiments – Enviroments

CFU (Custom Function Unit) Playground

System-on-Chip (SoC) equipped with a 32-bit RISC-V CPU + custom function unit

Tools

HDL: Chisel3

SW: C/C++

Outline

1. Abstract


2. Method

3. Experiments

- Environment
- **Serial vs SIMD**
- Summary


Experiments – Serial vs SIMD

Convolution




Experiments – Serial vs SIMD

GEMM



Experiments – Serial vs SIMD

Max Pooling



Outline

1. Abstract

2. Method

3. Experiments

- Environment
- Serial vs SIMD
- **Summary**

Experiments – Summary

	Serial	SIMD	SpeedUp
CONV	188M	123M	1.53
GEMM	746k	311k	2.4
MxPI	586k	290k	2.02

Thanks
